



**GMU Fall 2013, CS 780**  
**Class Project**

# **Partly Sunny with Chance of HashTags**

*a Kaggle Competition, Sponsored by CrowdFlower*

Presented by

**Talha Oz**  
**Venkat Tadakamalla**

# Agenda

1. Problem Description
2. Dataset
3. Goal
4. Tools and Software Packages Used
5. Processing Steps
6. Results

# 1. Problem Description

## Partly Sunny with Chance of HashTags

*a Kaggle Competition sponsored by*



**Training Data:** Approx. **78K** weather related tweets labeled by multiple people, crowd sourced !

**Test Data:** Approx. **42K**. Predict **24 classes** in 3 different categories.

- whether it has a positive, negative, or neutral sentiment
- whether the weather occurred in the past, present, or future
- what sort of weather the tweet references.


*Let us look at some more details...  
from Kaggle website*

# 1a. Problem Description, contd.

*Kaggle* specific requirements: No external info limits for this competition



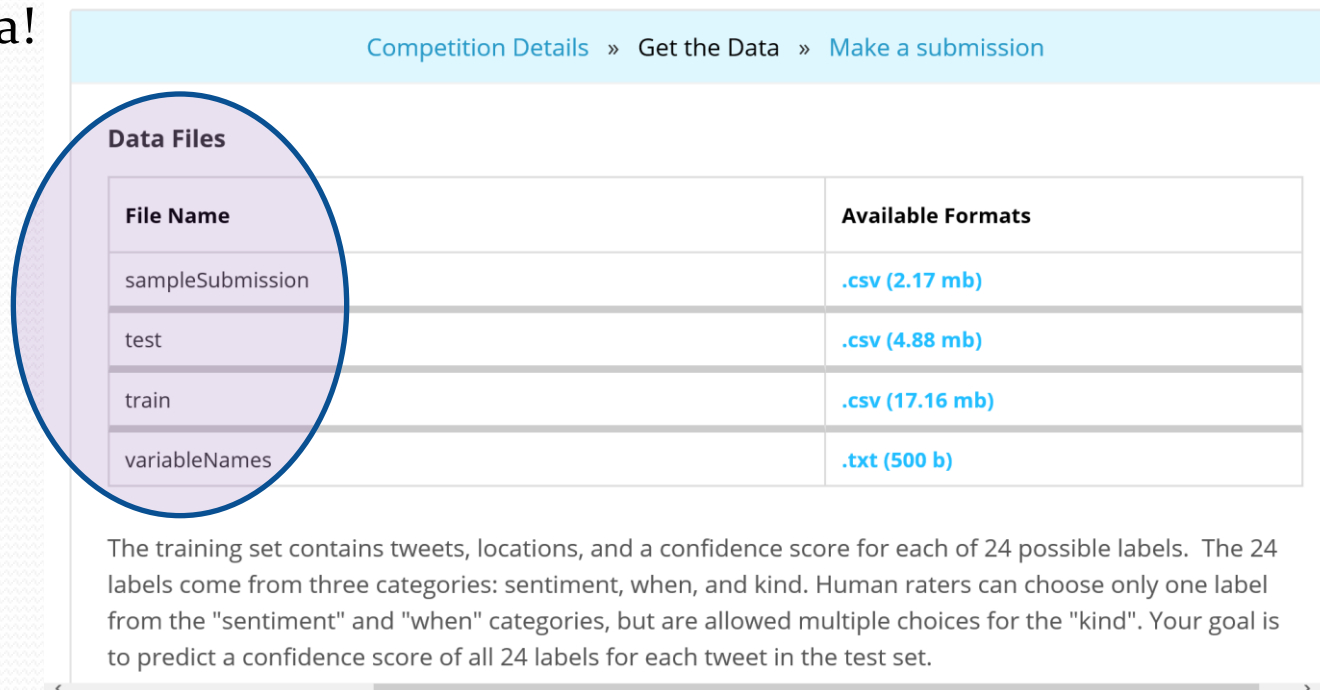
The screenshot shows the Kaggle website interface for the 'Partly Sunny with a Chance of Hashtags' competition. The top navigation bar includes links for Customer Solutions, Competitions, and Community. The competition page features a sidebar with navigation options like Dashboard, Home, Data, Make a submission, Information, Forum, Leaderboard, My Team, and My Submissions. The main content area displays the competition title, a progress bar indicating '4.2 days to go', and a description of the challenge: 'What can a #machine learn from tweets about the #weather?'. It also includes a leaderboard with the top 10 participants and a section for the competition details and rules.



**Started: 2:43 pm, Friday 27 September 2013 UTC**  
**Ends: 11:59 pm, Sunday 1 December 2013 UTC (65 total days)**

## 1b. Problem Description, contd.

- Compete with data miners all around the world (259 teams, 300 players)
- Daily upload limit for each team (5)
- Public/private leaderboard: 30% (public)/70% (private) of the test data!



Competition Details » Get the Data » Make a submission

Data Files	
File Name	Available Formats
sampleSubmission	<a href="#">.csv (2.17 mb)</a>
test	<a href="#">.csv (4.88 mb)</a>
train	<a href="#">.csv (17.16 mb)</a>
variableNames	<a href="#">.txt (500 b)</a>

The training set contains tweets, locations, and a confidence score for each of 24 possible labels. The 24 labels come from three categories: sentiment, when, and kind. Human raters can choose only one label from the "sentiment" and "when" categories, but are allowed multiple choices for the "kind". Your goal is to predict a confidence score of all 24 labels for each tweet in the test set.

## 1c. Problem Description, contd.

- 3 categories, **24 classes**:

Sentiment (5)	When (4)	Kind (15)
s1, "I can't tell"	w1, "current (same day) weather"	k1, "clouds"
s2, "Negative"	w2, "future (forecast)"	k2, "cold"
s3, "Neutral / author is just sharing information"	w3, "I can't tell"	k3, "dry"
s4, "Positive"	w4, "past weather"	k4, "hot"
s5, "Tweet not related to weather condition"		k5, "humid"
		k6, "hurricane"
		k7, "I can't tell"
		k8, "ice"
		k9, "other"
		k10, "rain"
		k11, "snow"
		k12, "storms"
		k13, "sun"
		k14, "tornado"
		k15, "wind"

## 1d. Problem Description, contd.

- Each tweet is labeled by multiple people
- Graders have reliability scores but we do not know them, instead:
- For When and Sentiment categories we have a probability distribution that sums up to one. Kind values can be more than one !

id	tweet	state	location
10	Rather be storm chasing.	district of columbia	Washington, DC
11	#WEATHER: 1:54 pm : 61.0F. Feels 60F. 29.98% Humidity. 6.9MPH East Wind.	michigan	St Louis, MI, USA

Y: S, W, K

(continuous variables)

s1	s2	s3	s4	s5	w1	w2	w3	w4
0	0	0.6	0	0.4	0	0	1	0
0	0	1	0	0	1	0	0	0

k1	k2	k3	k4	k5	k6	k7	k8	k9	k10	k11	k12	k13	k14	k15
0	0	0	0	0	0	0	0	0	0	0	1	0	0	0
0	0.203	0.176	0	0.376	0	0.421	0	0.176	0	0	0	0	0	0.579

✓ This is a multi-class continuous value problem !

## 2. Dataset

- Training Records (tweets) = 77,946
- Test Records (tweets) = 42,147

- Class Variables





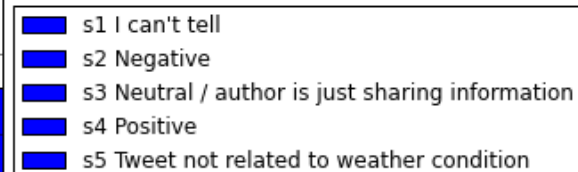

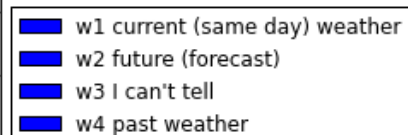

Sentiment	When	Kind
s1,"I can't tell"	w1,"current (same day) weather"	k1,"clouds"
s2,"Negative"	w2,"future (forecast)"	k2,"cold"
s3,"Neutral / author is just sharing information"	w3,"I can't tell"	k3,"dry"
s4,"Positive"	w4,"past weather"	k4,"hot"
s5,"Tweet not related to weather condition"		k5,"humid"
		k6,"hurricane"
		k7,"I can't tell"
		k8,"ice"
		k9,"other"
		k10,"rain"
		k11,"snow"
		k12,"storms"
		k13,"sun"
		k14,"tornado"
		k15,"wind"

- Sample tweets:

```
"1","Jazz for a Rainy Afternoon: {link}","oklahoma","Oklahoma", //id, tweet and location
"0","0","1","0","0", //sentiment
"0.8","0","0.2","0", //when
"0","0","0","0","0","0","0","0","0","1","0","0","0","0" //kind
"30","Today was windy AF, I don't wanna go meet up w/ @mention and get my sis from school
:/","arizona","Phoenix (Prescott)",
"0","0.613","0.387","0","0", //sentiment
"0.793","0","0","0.207", //when
"0","0","0","0","0","0","0","0","0","0","0","0","0","0","1" //kind
```



## 2a. Dataset, contd.



### 3. Goal

- Minimize the root mean square error !

$$\text{RMSE} = \sqrt{\frac{\sum_{i=1}^n (p_i - a_i)^2}{n}}$$

where:

- $n$  is 24 times the total number of tweets
  - $p_i$  is the predicted confidence rating for a given label
  - $a_i$  is the actual confidence rating for a given label
- Python code snippet:  
`np.sqrt(np.sum(np.array(pred-y_true)**2)/(y_true.shape[0]*float(y_true.shape[1])))`

# 4. Tools and Software Packages Used

- Python

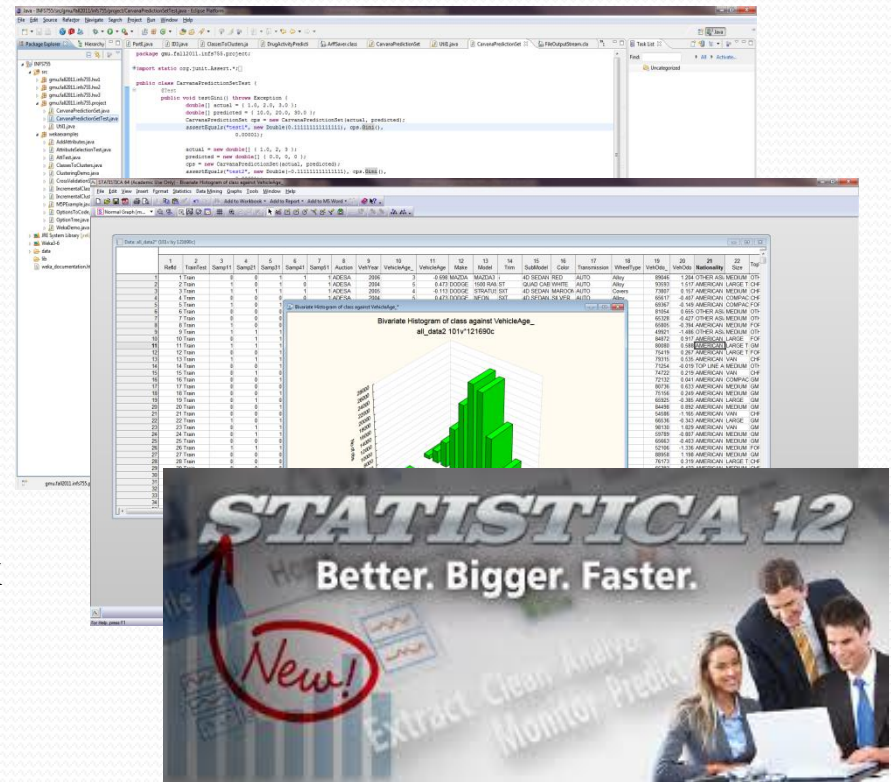
- Numpy, Scipy, Pandas
- Scikit-learn, NLTK

- Java

- Pre and Post Processing
- Explored MULAN, an open-source Java library based on WEKA for multi-class multi-label predictions. Found to work only for categorical values so discarded the tool.


- Statsoft's *STATISTICA*

- model building and predictions



## 5. Processing Steps

- a) Pre-process data
- b) Build models
- c) Evaluate models
- d) Predict and Deploy  
(upload to Kaggle site)



## 5a. Pre-process Data

- Feature Extraction/Creation from raw tweets
  - Raw words, Porter/Lancaster/Snowball stemmers, wordnet lemmatizer (NLTK)
  - Stop words removal (NLTK, scikit-learn)
  - Emoticons (happy/sad)
  - Unigrams, bigrams, together (scikit-learn)
- *STATISTICA*: for some selective runs, added additional features (Label\_1 to Label\_24) which are computed in case a synonym for any of the class label is found. Improved the RMSE by about 10%.

```
"cloud fog gloom smog smoke  
"cold chilly freezing bitter  
"dry moistureless arid bare  
"hot blazing boiling heated  
"humid damp dank muggy oppre  
"hurricane violent cyclone t  
"cant",  
"ice chunk crystal diamonds  
"other",  
"rain deluge drizzle flood r  
"snow blizzard snowfall",  
"storm squall",
```

## 5a. Pre-process Data, contd.

- Feature Selection/Reduction
  - Stacking state and location information to tweets
  - TFIDF vectorizer (max\_features, smoothing, min\_df, max\_df), count vectorizer
  - LSA (truncated SVD, work with sparse matrix, on sample vec)
  - PCA (dense only, on covariance matrix)
  - SelectKBest(chi2, class/feature dependency)
- *STATISTICA*: Utilized the best feature selection option for each individual classes. Used top 100 of such features for each class.

## 5b. Build Models

Used several algorithms

- **Regressors** (continuous, clip [0,1])
  - Ridge Regressor (sklearn)
  - SGD (Stochastic Gradient Descent) Regressor (sklearn)
- **Classifiers**
  - Predict\_proba
  - MultinomialNB (cs780 hw2, sklearn)
  - KNN
- **Ensemble methods**, Regression (Using *STATISTICA*\*)
  - Boosted Trees
  - Random Forests

*\* Noticed the following limitations (may or may not be accurate as we have a very limited time to explore the tool): Required the input data to be in sparse format; No multi-class support for many models. Typical test/ train file with about top 4000 features was approx. 1GB. Typical execution on a 64 bit AMD 2.5 GHz Quad Core PC with 8GB of RAM took about an hour (Execution of each model consisted of 24 individual executions, one for each class, manual assembling of results into CSV file).*

## 5c. Evaluate Models

- Cross Validation (everything in CV loop)
- Optimizing predictor variables (Grid Search)
- Optimizing token vectorizer variables



## 5d. Predict & Deploy

- With the best models figured out in the previous step, run the model on the Test Set to get the predictions.
- Upload the predictions to Kaggle web site.

The screenshot shows the Kaggle 'Make a submission' interface. On the left is a sidebar with a 'Dashboard' section containing links to Home, Data, Make a submission, Information, Forum, Leaderboard, My Team, and My Submissions. Below this is a 'Leaderboard' section listing the top 10 teams: 1. Tim Dettmers, 2. Duffman & Pascal, 3. David, 4. Jack Shih Tzu, 5. no\_name, 6. Deligne, 7. Razor, 8. LeastSquares, 9. YaTa, and 10. Anonymous.

The main content area is titled 'Make a submission' and includes a breadcrumb trail: 'Competition Details » Get the Data » Make a submission'. A message states: 'You have 5 entries today. This resets 3.3 hours from now (00:00 UTC)'. Below this is a large yellow box with the text 'Click or drop your submission here'. To the right of this box is a 'File Format' section stating: 'Your submission should be in CSV format. You can upload this in a zip/gz/rar/7z archive if you prefer.' Below the file format section is a '# of Predictions' section stating: 'We expect the solution file to have 42,157 predictions. The file should have a header row. Please see the sample submission file on the data page for an example of a valid submission.' At the bottom of the main content area is a blue 'Submit' button.

## 6. Results

- Huge part of the effort was **pre-processing** and **trying various different models** to see which model gives the best RMSE.
- Used Python and data mining libraries.
- Used **Custom Java programs** and *STATISTICA*
  - Developed custom Java programs for pre and post processing.
  - Normalized predicted class labels *STATISTICA* to sum to a desired value , e.g., 1.0 for sentiment and when labels.
- Built several models.
  - Models based on **Ridge Regressor** have performed extremely well.
- Uploaded many predictions onto *Kaggle* website, about 35 (1% of all participant uploads).

## 6. Results, contd.

Used several algorithms – Benchmark(all zeros: RMSE ~ 0.31597)

- **Regressors** (continuous, clip [0,1])
  - Ridge Regressor (sklearn) (**RMSE ~ 0.1568** ← best of all models)
  - SGD (Stochastic Gradient Descent) Regressor (sklearn)
- **Classifiers** (nominal, RMSE ~ 0.26 )
  - Predict\_proba
  - MultinomialNB (cs78o hw2, sklearn)
  - KNN
- **Ensemble methods**, Regression (Using *STATISTICA*)
  - Boosted Trees (RMSE ~ 0.22)
  - Random Forests (RMSE ~ 0.21 and ~ 0.19<sup>\*\*</sup>)

<sup>\*\*</sup> *With new additional features introduced based on the synonyms for the class labels*

## 6. Results, contd.

What **we learned** from the *Kaggle* online discussion forum:

- Simple stacking: feed the output of a first level model to use it as features to the 2nd level model
- Add more features: POS tagging, sentiment dictionary, etc. (**1<sup>st</sup> ranked fellow** had **~1.9M features**)

## 6. Results, contd.

- Competition started about **two months ago**. Some teams had **100+** submissions.

**Started:** 2:43 pm, Friday 27 September 2013 UTC  
**Ended:** 11:59 pm, Sunday 1 December 2013 UTC(65 total days)

- Final Numbers from *Kaggle* website:
- We entered **just two weeks ago** and are glad to report that we made it to the top **25%** of the leader board by the time the competition finished on Dec. 1<sup>st</sup> 2013.

259 teams  
300 players  
3604 entries



## References

- *NLTK (Natural Language Toolkit)*, <http://nltk.org/>
- *Scikit-learn Machine Learning in Python*, <http://scikit-learn.org/stable/>
- *Pandas Python Data Analysis Library*, <http://pandas.pydata.org/>
- *StatSoft 's STATISTICA 12 (64-bit edition) Academic License*, <http://www.statsoft.com/>



Thank you!  
&  
Questions?